

Threading Matrix Multiply

61519322 Yang Zherui

Date: 2021-5-15

Threading Matrix Multiply

Target 1: Matrix Multiply

Matrix Multiplication Function For Each Thread

Assign last few threads with extra rows

Moreover: Bind thread to specific CPU core

Target 2: Command Line Arguments and Timer

Command Line

Timer

Experiments:

Setup

Discussion about the result

Codes and Logs

main.c

log.txt

Targets:

1. A `matmul` using multiple thread acceleration.($C = A \cdot B$, we divide up rows of A for different threads)
2. Assign last few extra rows to different rows instead of assigning these rows to the last row only.
3. Specific command line arguments
4. Solution should be timed. (for multiple thread). and estimate the time for I/O and matrix multiplication.

Target 1: Matrix Multiply

Matrix Multiplication Function For Each Thread

In this section, I developed a matrix multiplication with low cache miss rate:

```
void * mm(void* arg)
{
    ...
    /* extract A, B, C, n, m, k from arg */
    for (int i = 0; i < n; ++i)
        for (int k_ = 0; k_ < m; ++k_)
            for (int j = 0; j < k; ++j)
                C[i * k + j] += A[i * m + k_] * B[k_ * k + j];
}
```

The matrix multiplication always happens here. To divide up the rows of input matrix A and the output matrix C , we can give a bias to A and C when using multiple thread: if we use `start` to represent the number of first line for each thread, and `end` to represent the number of last line for each thread, we have:

1. $A = \text{matA} + \text{start} * m$
2. $B = \text{matB} + \text{start} * m$

Secondly, when it comes to passing arguments to each thread, a pointer (`void*`) should be passed to each thread. To pass multiple arguments to each thread (for function `mm`), a struct as follow should be implemented:

```
struct mm_info {
    float *A, *B, *out;
    int n, m, k;
};
```

In this case, we can implement the function `mm` as follows:

```
void *mm(void *info)
{
    /* 提取需要的信息 */
    struct mm_info *pinfo = (struct mm_info *)info;
    float *A = pinfo->A;
    float *B = pinfo->B;
    int n = pinfo->n,
```

```

        m = pinfo->m,
        k = pinfo->k;
float *C = pinfo->out;
/* 清空该线程需要的 C (子) 矩阵 */
for (int i = 0; i < n * k; ++i)
    C[i] = 0.0f;

for (int i = 0; i < n; ++i)
    for (int k_ = 0; k_ < m; ++k_)
        for (int j = 0; j < k; ++j)
            C[i * k + j] += A[i * m + k_] * B[k_ * k + j];
}

```

Assign last few threads with extra rows

In this section, use:

1. `nthread` to denote the number of threads we use in the progress.
2. `rows_each_thread` to denote the number of rows for each normal threads to do `mm()`.
3. `extra_rows_to_assign` to denote the number of threads to do `mm()` with extra lines.
4. `n` to denote the number of rows of matrix *A*
5. `it` to denote the rank of each thread

To obtain a better performance for the whole progress, we can use:

1. `int rows_each_thread = n / nthread;`
2. `int extra_rows_to_assign = nthread - n / rows_each_thread * nthread;`

Therefore, to generate each thread:

```

/* 存储各个线程矩阵乘法需要的信息 */
struct mm_info *t_info = malloc(sizeof(struct mm_info) * nthread);
/* 存储各个线程的 tid */
pthread_t *tid = malloc(sizeof(pthread_t) * nthread);
for (int it = 0; it < nthread; ++it){
    int start = rows_each_thread * it;
    if (it >= extra_rows_to_assign)
        start += it - extra_rows_to_assign;
    int end = start + rows_each_thread;
    if (it >= extra_rows_to_assign)
        end += 1;
}

```

```

if (end > n_A) end = n_A;
/* 填写线程参数 */
t_info[it].A = A + start * m_A;
t_info[it].B = B;
t_info[it].n = end - start;
t_info[it].m = m_A;
t_info[it].k = m_B;
t_info[it].out = C + start * m_C;
/* 创建线程 */
pthread_create(&tid[it], NULL, mm, t_info + it);
}

```

Moreover: Bind thread to specific CPU core

To maximize the performance, we can manually bind threads to specific CPU cores by:

```

cpu_set_t cpu;
CPU_ZERO(&cpu);
CPU_SET(it % 8, &cpu);
...
pthread_setaffinity_np(tid[it], sizeof(cpu_set_t), &cpu);

```

Then, each thread will be allocated averagely to each CPU core.

Target 2: Command Line Arguments and Timer

Command Line

In this program, there is only one format of command line arguments can be passed.

```
./mm -a Amat.txt -b Bmat.txt -t 4
```

Therefore, the preprocess section can be written as follows:

```
int main(int argc, char **argv)
```

```

{
    if (argc != 7)
    {
        printf("[Error] Illegal Param For matmul.");
        return 1;
    }

    int m_A, m_B, n_A, n_B;
    float *A, *B, *C;
    float start,
          mm_start,
          end;
    A = read_matrix(&m_A, &n_A, argv[2]);
    B = read_matrix(&m_B, &n_B, argv[4]);
    int nthread = atoi(argv[6]);
    int m_C = m_B,
          n_C = n_A;
    C = (float *)malloc(sizeof(float) * n_C * m_C);

    /* Do mm ...*/

    /* Free mm spaces */
    return 0;
}

```

Timer

*There is **always a bunch of problems** when it comes to implement a timer when using multiple thread programming!*

There is several ways to get the running time of a program:

1. `time(NULL)` (in `time.h`) get the time measured in seconds: the output for a multiple thread program is just the same as a single thread program;
2. `clock()` (in `time.h`) get the CPU clock of this program, divide `CLOCKS_PER_SEC` to get the CPU time elapsed during the program. However, this will make a difference under the environment of multi-thread programming. This function will return the TOTAL clock of this progress, including every thread's clock ticks.
3. `clock_gettime(clockid_t clk_id, struct timespec *res)` : finds the resolution (precision) of the specified clock `clk_id` , and, if `res` is non-NULL, stores it in the struct `timespec` pointed to by `res`. If we use `clk_id = CLOCK_MONOTONIC` , we can

get the precise time corresponds to the number of seconds that the system has been running since it was booted.

Therefore, we can call `clock_gettime(CLOCK_MONOTONIC, &t)` to get the number of seconds elapsed during my program, regardless that the program is running under a multi-thread environment.

```
#define TIME_MS(t) (((float) t.tv_sec) * 1000.f + ((float) t.tv_nsec) /  
1e6f)  
...  
// in main function  
struct timespec t;  
clock_gettime(CLOCK_MONOTONIC, &t);  
start = TIME_MS(t);  
...  

```

Experiments:

Setup

To run the whole test, I use a shell script as follows:

```
gcc main.c -O2 -lpthread -pthread -o mm  
for size in 2 4 8 10 100 200 500 1000 2000 4000 8000  
do  
    echo "#####" $size "#####"  
    python gm.py $size $size "Amat$size.txt"  
    python gm.py $size $size "Bmat$size.txt"  
    for nt in 1 2 4 8 16  
    do  
        echo "mm -a Amat$size.txt -b Bmat$size.txt -t $nt"  
  
        time -p ./mm -a Amat$size.txt -b Bmat$size.txt -t $nt  
    done  
done  
echo "[DONE.]"
```

and run it using the following command to get the whole log.


```
$ sh test.sh > log.txt
```

Discussion about the result

- t\size	2	4	8	10	100	200	500	1000	2000	4000	8000
2	0.5	0.5	0.5	5	13.5	44.5	302	1428	8395	50825	353577
4	1	1	1.5	1	14	44.5	286	1184	6355	34971	222900
6	1	1.5	2	1	14	47.5	254	1069	5214	28159	173837
8	1.5	1.5	1.5	1.5	16	41.5	242	1044	4841	24701	157646
16	1.5	1.5	2	1.5	18	48	252	1191	4892	26052	161803

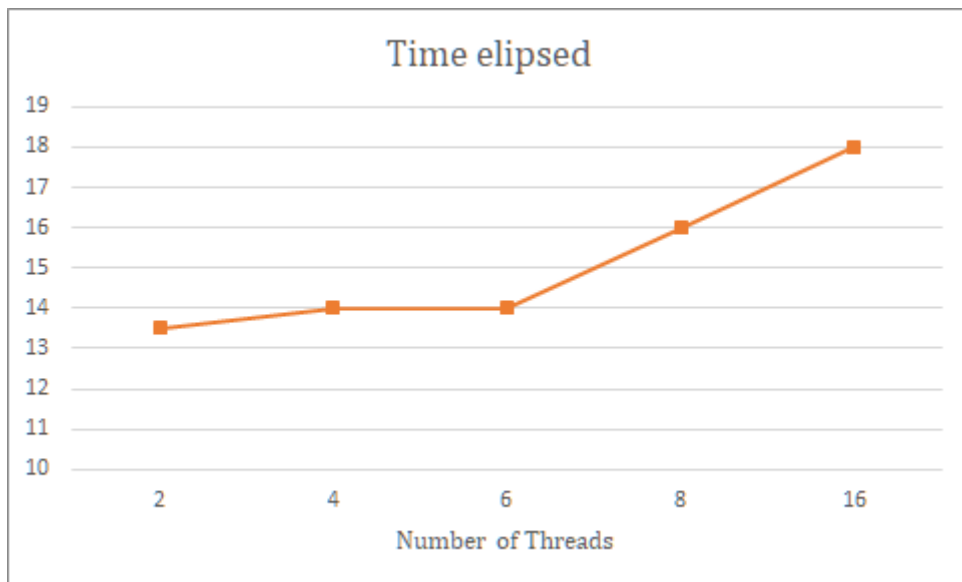
Table: Time elapsed (in ms) using different arguments for the program (-t = number of threads, size = the size of matrices) Tested on a 4-core 8-thread CPU

Take the size = 8000 as an example :



It is obvious that, with multiple threads, there will be a obvious acceleration for matrix multiplication. But when the thread is greater that the number of CPU cores, the performance will be lower than a small thread number.

Another example is size = 100:



With a low size but a high number of threads, more time will be spent on creation and joining of threads.

Codes and Logs

1. `main.c` matrix multiplication program
2. `log.txt` test log.

main.c

```
#define _GNU_SOURCE
#include <stdlib.h>
#include <pthread.h>
#include <stdio.h>
#include <string.h>
#include <time.h>
#include <sched.h>
#define TIME_MS(t) (((float) t.tv_sec) * 1000.f + ((float) t.tv_nsec) / 1e6f)

#define NDEBUG

struct mm_info
{
    float *A;
```



```

    float *B;
    int n, m, k;
    float *out;
};

// do matmul store in out
void *mm(void *info)
{
    struct mm_info *pinfo = (struct mm_info *)info;
    float *A = pinfo->A;
    float *B = pinfo->B;
    int n = pinfo->n,
        m = pinfo->m,
        k = pinfo->k;
    float *C = pinfo->out;

#ifdef NDEBUG
    printf("[Info] I'm going to do matmul...\n");
    printf("\tMat A: address = %p\n", A);
    printf("\tMat B: address = %p\n", B);
    printf("\tMat C: address = %p\n", C);
    printf("\tn, m, k = %d, %d, %d\n", n, m, k);
#endif

    for (int i = 0; i < n * k; ++i)
        C[i] = 0.0f;

    for (int i = 0; i < n; ++i)
    {
        for (int k_ = 0; k_ < m; ++k_)
        {
            for (int j = 0; j < k; ++j)
            {
                C[i * k + j] += A[i * m + k_] * B[k_ * k + j];
            }
        }
    }
}

// read from text
float *read_matrix(int *m, int *n, char *fn)
{
    FILE *file = fopen(fn, "r");
    if (file == NULL)

```

```

    {
        printf("[Error] I cannot read file %s.\n", fn);
    }
    fscanf(file, "%d", n);
    fscanf(file, "%d", m);
    float *mat = (float *)malloc(sizeof(float) * (*n) * (*m));
    for (int i = 0; i < (*n) * (*m); ++i)
    {
        fscanf(file, "%f", &(mat[i]));
    }

    fclose(file);
    return mat;
}

// write to text
void write_matrix(float *mat, int *m, int *n, char *fn)
{
    FILE *file = fopen(fn, "w");
    fprintf(file, "%d %d\n", *m, *n);
    for (int i = 0; i < *n; ++i)
    {
        for (int j = 0; j < *m; ++j)
        {
            fprintf(file, "%f ", mat[i * (*m) + j]);
        }
        fprintf(file, "\n");
    }
    fclose(file);
}

int main(int argc, char **argv)
{
    // mm -a mata.txt -b matb.txt -t nthreads
    if (argc != 7)
    {
        printf("[Error] Illegal Param For matmul.");
        return 1;
    }

    int m_A, m_B, n_A, n_B;
    float *A, *B, *C;
    float start,

```

```

        mm_start,
        end;

    struct timespec t;
    clock_gettime(CLOCK_MONOTONIC, &t);
    start = TIME_MS(t);

#ifdef NDEBUG
    printf("[Info] I wanna read A from %s and B from %s...\n", argv[2],
argv[4]);
#endif

    A = read_matrix(&m_A, &n_A, argv[2]);
    B = read_matrix(&m_B, &n_B, argv[4]);
#ifdef NDEBUG
    printf("[Info] I got mat A <at %p> and mat B <at %p>\n", A, B);
#endif

    int nthread = atoi(argv[6]);
    if (m_A != n_B)
    {
        printf("[Error] Illegal shape For matmul, got m1, m2 = [%d, %d]",
m_A, n_B);
    }

    int m_C = m_B,
        n_C = n_A;
    C = (float *)malloc(sizeof(float) * n_C * m_C);

    int rows_each_thread = n_A / nthread;
    int extra_rows_to_assign = nthread - n_A + rows_each_thread *
nthread;

    struct mm_info *t_info = malloc(sizeof(struct mm_info) * nthread);
    pthread_t *tid = malloc(sizeof(pthread_t) * nthread);
    clock_gettime(CLOCK_MONOTONIC, &t);
    mm_start = TIME_MS(t);
    printf("[Info] Done IO in %lf ms\n", (((double) (mm_start -
start)))));
    for (int it = 0; it < nthread; ++it)
    {
        cpu_set_t cpu;

```

```

CPU_ZERO(&cpu);
CPU_SET(it % 8, &cpu);
int start = rows_each_thread * it;
if (it >= extra_rows_to_assign)
    start += it - extra_rows_to_assign;
int end = start + rows_each_thread;
if (it >= extra_rows_to_assign)
    end += 1;
if (end > n_A) end = n_A;
t_info[it].A = A + start * m_A;
t_info[it].B = B;
t_info[it].n = end - start;
t_info[it].m = m_A;
t_info[it].k = m_B;
t_info[it].out = C + start * m_C;
#ifdef NDEBUG
    printf("n, m, k = %d, %d, %d\n", t_info[it].n, t_info[it].m,
t_info[it].k);
#endif
    pthread_create(tid + it, NULL, mm, t_info + it);
    pthread_setaffinity_np(tid[it], sizeof(cpu_set_t), &cpu);
}

for (int it = 0; it < nthread; ++it){
    pthread_join(tid[it], NULL);
}
clock_gettime(CLOCK_MONOTONIC, &t);
end = TIME_MS(t);
printf("[Info] Done MM in %lf ms...\n", (((double) (end -
mm_start)))));

write_matrix(C, &m_C, &n_C, "result.txt");

clock_gettime(CLOCK_MONOTONIC, &t);
end = TIME_MS(t);
printf("[Info] Done in %lf ms...\n", ((double) (end - start)));

free(A); A = NULL;
free(B); B = NULL;
free(C); C = NULL;
free(tid); tid = NULL;
free(t_info); t_info = NULL;

```

```
    return 0;
}
```

log.txt

```
##### 2 #####
mm -a Amat2.txt -b Bmat2.txt -t 1
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 5.000000 ms...
mm -a Amat2.txt -b Bmat2.txt -t 2
[Info] Done IO in 0.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.000000 ms...
mm -a Amat2.txt -b Bmat2.txt -t 4
[Info] Done IO in 0.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.000000 ms...
mm -a Amat2.txt -b Bmat2.txt -t 8
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.500000 ms...
mm -a Amat2.txt -b Bmat2.txt -t 16
[Info] Done IO in 0.000000 ms
[Info] Done MM in 1.000000 ms...
[Info] Done in 1.500000 ms...
##### 4 #####
mm -a Amat4.txt -b Bmat4.txt -t 1
[Info] Done IO in 0.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 4.500000 ms...
mm -a Amat4.txt -b Bmat4.txt -t 2
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 1.000000 ms...
mm -a Amat4.txt -b Bmat4.txt -t 4
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 1.500000 ms...
```

```
mm -a Amat4.txt -b Bmat4.txt -t 8
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.500000 ms...
mm -a Amat4.txt -b Bmat4.txt -t 16
[Info] Done IO in 0.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.500000 ms...
##### 8 #####
mm -a Amat8.txt -b Bmat8.txt -t 1
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 5.000000 ms...
mm -a Amat8.txt -b Bmat8.txt -t 2
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 1.500000 ms...
mm -a Amat8.txt -b Bmat8.txt -t 4
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 2.000000 ms...
mm -a Amat8.txt -b Bmat8.txt -t 8
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.500000 ms...
mm -a Amat8.txt -b Bmat8.txt -t 16
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 2.000000 ms...
##### 10 #####
mm -a Amat10.txt -b Bmat10.txt -t 1
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 5.000000 ms...
mm -a Amat10.txt -b Bmat10.txt -t 2
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 1.000000 ms...
mm -a Amat10.txt -b Bmat10.txt -t 4
[Info] Done IO in 0.500000 ms
[Info] Done MM in 0.000000 ms...
[Info] Done in 1.000000 ms...
mm -a Amat10.txt -b Bmat10.txt -t 8
```

```
[Info] Done IO in 0.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 1.500000 ms...
mm -a Amat10.txt -b Bmat10.txt -t 16
[Info] Done IO in 0.000000 ms
[Info] Done MM in 1.000000 ms...
[Info] Done in 1.500000 ms...
##### 100 #####
mm -a Amat100.txt -b Bmat100.txt -t 1
[Info] Done IO in 4.500000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 13.500000 ms...
mm -a Amat100.txt -b Bmat100.txt -t 2
[Info] Done IO in 5.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 14.000000 ms...
mm -a Amat100.txt -b Bmat100.txt -t 4
[Info] Done IO in 4.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 14.000000 ms...
mm -a Amat100.txt -b Bmat100.txt -t 8
[Info] Done IO in 5.000000 ms
[Info] Done MM in 0.500000 ms...
[Info] Done in 16.000000 ms...
mm -a Amat100.txt -b Bmat100.txt -t 16
[Info] Done IO in 5.000000 ms
[Info] Done MM in 1.000000 ms...
[Info] Done in 18.000000 ms...
##### 200 #####
mm -a Amat200.txt -b Bmat200.txt -t 1
[Info] Done IO in 17.000000 ms
[Info] Done MM in 5.000000 ms...
[Info] Done in 44.500000 ms...
mm -a Amat200.txt -b Bmat200.txt -t 2
[Info] Done IO in 17.500000 ms
[Info] Done MM in 2.500000 ms...
[Info] Done in 44.500000 ms...
mm -a Amat200.txt -b Bmat200.txt -t 4
[Info] Done IO in 19.500000 ms
[Info] Done MM in 2.000000 ms...
[Info] Done in 47.500000 ms...
mm -a Amat200.txt -b Bmat200.txt -t 8
[Info] Done IO in 17.000000 ms
```

```
[Info] Done MM in 2.000000 ms...
[Info] Done in 41.500000 ms...
mm -a Amat200.txt -b Bmat200.txt -t 16
[Info] Done IO in 20.500000 ms
[Info] Done MM in 2.000000 ms...
[Info] Done in 48.000000 ms...
##### 500 #####
mm -a Amat500.txt -b Bmat500.txt -t 1
[Info] Done IO in 114.000000 ms
[Info] Done MM in 70.000000 ms...
[Info] Done in 302.000000 ms...
mm -a Amat500.txt -b Bmat500.txt -t 2
[Info] Done IO in 117.500000 ms
[Info] Done MM in 48.000000 ms...
[Info] Done in 286.000000 ms...
mm -a Amat500.txt -b Bmat500.txt -t 4
[Info] Done IO in 107.500000 ms
[Info] Done MM in 26.000000 ms...
[Info] Done in 254.500000 ms...
mm -a Amat500.txt -b Bmat500.txt -t 8
[Info] Done IO in 106.500000 ms
[Info] Done MM in 18.500000 ms...
[Info] Done in 242.000000 ms...
mm -a Amat500.txt -b Bmat500.txt -t 16
[Info] Done IO in 116.000000 ms
[Info] Done MM in 18.500000 ms...
[Info] Done in 252.500000 ms...
##### 1000 #####
mm -a Amat1000.txt -b Bmat1000.txt -t 1
[Info] Done IO in 440.500000 ms
[Info] Done MM in 524.000000 ms...
[Info] Done in 1428.000000 ms...
mm -a Amat1000.txt -b Bmat1000.txt -t 2
[Info] Done IO in 421.500000 ms
[Info] Done MM in 292.000000 ms...
[Info] Done in 1184.000000 ms...
mm -a Amat1000.txt -b Bmat1000.txt -t 4
[Info] Done IO in 421.000000 ms
[Info] Done MM in 187.000000 ms...
[Info] Done in 1069.500000 ms...
mm -a Amat1000.txt -b Bmat1000.txt -t 8
[Info] Done IO in 430.500000 ms
[Info] Done MM in 148.500000 ms...
```



```
[Info] Done in 1044.000000 ms...
mm -a Amat1000.txt -b Bmat1000.txt -t 16
[Info] Done IO in 445.500000 ms
[Info] Done MM in 169.500000 ms...
[Info] Done in 1191.000000 ms...
##### 2000 #####
mm -a Amat2000.txt -b Bmat2000.txt -t 1
[Info] Done IO in 1749.000000 ms
[Info] Done MM in 4744.000000 ms...
[Info] Done in 8395.000000 ms...
mm -a Amat2000.txt -b Bmat2000.txt -t 2
[Info] Done IO in 1690.500000 ms
[Info] Done MM in 2847.000000 ms...
[Info] Done in 6355.000000 ms...
mm -a Amat2000.txt -b Bmat2000.txt -t 4
[Info] Done IO in 1674.000000 ms
[Info] Done MM in 1606.500000 ms...
[Info] Done in 5214.500000 ms...
mm -a Amat2000.txt -b Bmat2000.txt -t 8
[Info] Done IO in 1702.000000 ms
[Info] Done MM in 1274.500000 ms...
[Info] Done in 4841.000000 ms...
mm -a Amat2000.txt -b Bmat2000.txt -t 16
[Info] Done IO in 1790.500000 ms
[Info] Done MM in 1276.500000 ms...
[Info] Done in 4862.000000 ms...
##### 4000 #####
mm -a Amat4000.txt -b Bmat4000.txt -t 1
[Info] Done IO in 6706.500000 ms
[Info] Done MM in 36953.500000 ms...
[Info] Done in 50825.500000 ms...
mm -a Amat4000.txt -b Bmat4000.txt -t 2
[Info] Done IO in 6621.000000 ms
[Info] Done MM in 21158.500000 ms...
[Info] Done in 34971.000000 ms...
mm -a Amat4000.txt -b Bmat4000.txt -t 4
[Info] Done IO in 6584.500000 ms
[Info] Done MM in 14384.500000 ms...
[Info] Done in 28159.500000 ms...
mm -a Amat4000.txt -b Bmat4000.txt -t 8
[Info] Done IO in 6634.000000 ms
[Info] Done MM in 10832.000000 ms...
[Info] Done in 24701.000000 ms...
```

```
mm -a Amat4000.txt -b Bmat4000.txt -t 16
[Info] Done IO in 6758.500000 ms
[Info] Done MM in 12039.500000 ms...
[Info] Done in 26052.000000 ms...
##### 8000 #####
mm -a Amat8000.txt -b Bmat8000.txt -t 1
[Info] Done IO in 30754.000000 ms
[Info] Done MM in 292763.000000 ms...
[Info] Done in 353577.500000 ms...
mm -a Amat8000.txt -b Bmat8000.txt -t 2
[Info] Done IO in 26304.000000 ms
[Info] Done MM in 166687.000000 ms...
[Info] Done in 222900.000000 ms...
mm -a Amat8000.txt -b Bmat8000.txt -t 4
[Info] Done IO in 26186.000000 ms
[Info] Done MM in 117384.000000 ms...
[Info] Done in 173837.500000 ms...
mm -a Amat8000.txt -b Bmat8000.txt -t 8
[Info] Done IO in 26474.500000 ms
[Info] Done MM in 100375.000000 ms...
[Info] Done in 157646.500000 ms...
mm -a Amat8000.txt -b Bmat8000.txt -t 16
[Info] Done IO in 26634.500000 ms
[Info] Done MM in 104828.000000 ms...
[Info] Done in 161803.500000 ms...
[DONE.]
```